

Congressional Voting First Model

March 30, 2021

1 Congressional Voting First Model Attempt

1.1 Bartosz Sawicki, Hubert Ruczyński

```
[1]: import numpy as np
import pandas as pd
import requests
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

from matplotlib import pyplot as plt
plt.style.use('ggplot')

from sklearn import metrics
from sklearn import tree
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import cross_validate
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
from sklearn.pipeline import Pipeline

from xgboost import XGBClassifier # Inna paczka niż sklearn!
```

2 Wczytanie Danych

```
[2]: url = 'https://api.apispreadsheets.com/api/dataset/congressional-voting/'

r = requests.get(url)
data = r.json()
df = pd.DataFrame.from_dict(data['data'], orient='columns')
df.sample(5)
```

```
[2]:      handicapped_infants  water_project_cost_sharing  \
119                        n                             n
342                        n                             y
```

24	y	n
405	n	n
290	y	n

	adoption_of_the_budget_resolution	physician_fee_freeze	el_salvador_aid	\
119	n	y	y	
342	y	n	y	
24	y	n	n	
405	n	y	y	
290	y	n	?	

	religious_groups_in_schools	anti_satellite_test_ban	\
119	y	n	
342	?	y	
24	n	y	
405	y	n	
290	y	?	

	aid_to_nicaraguan_contras	mx_missile	immigration	\
119	n	n	n	
342	n	n	y	
24	y	y	n	
405	n	n	n	
290	y	y	y	

	synfuels_corporation_cutback	education_spending	superfund_right_to_sue	\
119	n	y	y	
342	y	n	y	
24	n	n	n	
405	n	y	y	
290	n	n	y	

	crime	duty_free_exports	export_administration_act_south_africa	\
119	y	n	n	
342	n	y	y	
24	n	y	?	
405	y	n	y	
290	y	n	y	

	political_party
119	republican
342	democrat
24	democrat
405	republican
290	democrat

3 Autorski Encoding

```
[3]: val = {'n':-1, '?':0, 'y':1}
party = {'democrat':2, 'republican':-2}

df_num=df.copy()
df_party = df_num['political_party'].copy()
df_num.drop('political_party', axis=1, inplace=True)

for column in df_num.columns:
    df_num[column] = df_num[column].map(val)

df_party = df_party.map(party)

df_num = pd.concat([df_num,df_party], axis=1)
df_num.head()
```

```
[3]:   handicapped_infants  water_project_cost_sharing  \
0                      -1                        1
1                      -1                        1
2                       0                        1
3                      -1                        1
4                       1                        1

      adoption_of_the_budget_resolution  physician_fee_freeze  el_salvador_aid  \
0                                -1                        1            1
1                                -1                        1            1
2                                 1                        0            1
3                                 1                       -1            0
4                                 1                       -1            1

      religious_groups_in_schools  anti_satellite_test_ban  \
0                                1                        -1
1                                1                        -1
2                                1                        -1
3                                1                        -1
4                                1                        -1

      aid_to_nicaraguan_contras  mx_missile  immigration  \
0                                -1          -1           1
1                                -1          -1          -1
2                                -1          -1          -1
3                                -1          -1          -1
4                                -1          -1          -1

      synfuels_corporation_cutback  education_spending  superfund_right_to_sue  \
0                                0                      1                      1
```

1		-1	1	1
2		1	-1	1
3		1	-1	1
4		1	0	1

	crime	duty_free_exports	export_administration_act_south_africa	\
0	1	-1	1	
1	1	-1	0	
2	1	-1	-1	
3	-1	-1	1	
4	1	1	1	

	political_party
0	-2
1	-2
2	2
3	2
4	2

4 Informacje o Danych

```
[4]: x = df_num.drop('political_party', axis=1)
y = df_num.iloc[:,16]
x.info()
y
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 435 entries, 0 to 434
```

```
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	handicapped_infants	435 non-null	int64
1	water_project_cost_sharing	435 non-null	int64
2	adoption_of_the_budget_resolution	435 non-null	int64
3	physician_fee_freeze	435 non-null	int64
4	el_salvador_aid	435 non-null	int64
5	religious_groups_in_schools	435 non-null	int64
6	anti_satellite_test_ban	435 non-null	int64
7	aid_to_nicaraguan_contras	435 non-null	int64
8	mx_missile	435 non-null	int64
9	immigration	435 non-null	int64
10	synfuels_corporation_cutback	435 non-null	int64
11	education_spending	435 non-null	int64
12	superfund_right_to_sue	435 non-null	int64
13	crime	435 non-null	int64
14	duty_free_exports	435 non-null	int64

```
15 export_administration_act_south_africa 435 non-null int64
dtypes: int64(16)
memory usage: 54.5 KB
```

```
[4]: 0      -2
      1      -2
      2       2
      3       2
      4       2
      ..
     430     -2
     431       2
     432     -2
     433     -2
     434     -2
      Name: political_party, Length: 435, dtype: int64
```

5 Pierwszy model: SVM

Model: <https://scikit-learn.org/stable/modules/svm.html#classification>

Za pierwszym razem nie dokonaliśmy selekcji zmiennych i z tego powodu otrzymaliśmy dość nieskuteczny model.

Accuracy powyżej 0,9 może być mylące gdyż na pierwszy rzut wydaje się, że jest to wynik bardzo dobry, jednak po analizie której dokonywaliśmy ostatnio, wiemy, że pojedyncze głosowanie (physicians_fee) jest w stanie przewidywać z dokładnością prawie 94%.

```
[5]: x_train, x_test, y_train, y_test = train_test_split(x, y)
```

```
[6]: from sklearn import svm
      clf = svm.SVC()
      clf.fit(x_train, y_train)
      y_train_pred=clf.predict(x_train)
```

```
[7]: y_pred=clf.predict(x_test)
```

```
[8]: print('\ny:      ' + str(y_test[0:10]) + '\ny_pred: ' + str(y_pred[0:10]))
      sum=0
      y_test=y_test.to_numpy()
      for i in range(len(y_pred)):
          if y_pred[i]==y_test[i]:
              sum+=1
      accuracy=sum/len(y_pred)
      print("validation : " + str(accuracy))
      sum=0
      y_trained=clf.predict(x_train)
      y_trains=y_train.to_numpy()
```

```

for i in range(len(y_trained)):
    if y_trained[i]==y_trains[i]:
        sum+=1
accuracy_train=sum/len(y_trained)
print("train : " + str(accuracy_train))

```

```

y:      59      -2
72       2
429      2
153      2
282     -2
143      2
402     -2
198      2
103      2
23       2
Name: political_party, dtype: int64
y_pred: [-2  2  2  2 -2  2 -2  2  2  2]
validation : 0.944954128440367
train : 0.9846625766871165

```

Aby sprawdzić faktyczne accuracy, postanowiliśmy zapętlić wybór zbioru testowego, tworzenia i trenowania modelu 1000 razy, aby otrzymać średnie accuracy dla tych prób.

Ponadto wykonaliśmy selekcji zmiennych przy tym testowaniu w dwóch wariantach i okazało się, że lepiej wypadł ten przy mniejszej liczbie wyrzucanych kolumn.

5.1 Uwaga!

Ilość testowanych modeli i zbiory wyrzucanych kolumn były znacznie większe w fazie pisania raportu, jednak postanowiliśmy zostawić tylko te które coś wnoszą.

```

[9]: accumulated_accuracy=0
accumulated_accuracy_train=0
for i in range(1000):
    x = df_num.drop(['political_party',
    ↪ 'water_project_cost_sharing', 'immigration', 'export_administration_act_south_africa'],
    ↪ axis=1)
    y = df_num.iloc[:,16]
    x_train, x_test, y_train, y_test = train_test_split(x, y)
    clf = svm.SVC()
    clf.fit(x_train, y_train)
    y_train_pred=clf.predict(x_train)
    y_pred=clf.predict(x_test)
    sum=0

    y_test=y_test.to_numpy()
    for i in range(len(y_pred)):

```

```

        if y_pred[i]==y_test[i]:
            sum+=1
    accuracy=sum/len(y_pred)
    accumulated_accuracy=accumulated_accuracy+accuracy
    sum=0
    y_train=clf.predict(x_train)
    y_train=y_train.to_numpy()
    for i in range(len(y_train)):
        if y_train[i]==y_train[i]:
            sum+=1
    accuracy_train=sum/len(y_train)
    accumulated_accuracy_train=accumulated_accuracy_train+accuracy_train

print("accumulated accuracy validation : " + str(accumulated_accuracy/1000))
print("accumulated accuracy train : " + str(accumulated_accuracy_train/1000))

```

accumulated accuracy validation : 0.9565229357798124
 accumulated accuracy train : 0.9793343558282152

```

[10]: accumulated_accuracy=0
    accumulated_accuracy_train=0
    for i in range(1000):
        x = df_num.drop(['political_party',
        ↳ 'water_project_cost_sharing', 'immigration', 'export_administration_act_south_africa',
        ↳ 'education_spending',
        ↳ 'synfuels_corporation_cutback', 'religious_groups_in_schools'], axis=1)
        y = df_num.iloc[:,16]
        x_train, x_test, y_train, y_test = train_test_split(x, y)
        clf = svm.SVC()
        clf.fit(x_train, y_train)
        y_train_pred=clf.predict(x_train)
        y_pred=clf.predict(x_test)
        sum=0

        y_test=y_test.to_numpy()
        for i in range(len(y_pred)):
            if y_pred[i]==y_test[i]:
                sum+=1
        accuracy=sum/len(y_pred)
        accumulated_accuracy=accumulated_accuracy+accuracy
        sum=0
        y_train=clf.predict(x_train)
        y_train=y_train.to_numpy()
        for i in range(len(y_train)):
            if y_train[i]==y_train[i]:
                sum+=1
        accuracy_train=sum/len(y_train)
    
```

```

        accumulated_accuracy_train=accumulated_accuracy_train+accuracy_train

print("accumulated accuracy validation : " + str(accumulated_accuracy/1000))
print("accumulated accuracy train : " + str(accumulated_accuracy_train/1000))

```

```

accumulated accuracy validation : 0.9484403669724762
accumulated accuracy train : 0.9693190184049109

```

Na autorskim kodowaniu postanowiliśmy też wykorzystać bardziej złożony model w postaci XGB-Classifier'a, który zapewnił nam o wiele wyższe accuracy od poprzedników.

```

[11]: x = df_num.drop(['political_party',
    ↪ 'water_project_cost_sharing', 'immigration', 'export_administration_act_south_africa'],
    ↪ axis=1)
y = df_num.iloc[:,16]

```

```

[12]: x_train, x_test, y_train, y_test = train_test_split(x, y, random_state=14)

```

```

[13]: model=XGBClassifier(random_state=2,
    ↪ learning_rate=0.01, # Szybkość "uczenia" się
    ↪ booster='gbtree', # Jaki model wykorzystujemy (drzewo -
    ↪ ↪ gbtree, liniowe - gblinear)
    ↪ nround = 1000, # Ilość iteracji boostingowych
    ↪ max_depth=3, # Maksymalna głębokość drzewa
    ↪ verbosity = 0 # Nie chcemy warningów
    ↪ )
model.fit(x_train, y_train)
prediction_test=model.predict(x_test)
print(model.score(x_test,y_test))
print(metrics.classification_report(y_test, prediction_test))

```

```

0.963302752293578

```

	precision	recall	f1-score	support
-2	0.94	0.98	0.96	46
2	0.98	0.95	0.97	63
accuracy			0.96	109
macro avg	0.96	0.97	0.96	109
weighted avg	0.96	0.96	0.96	109

6 Kodowanie one hot

Aby uzyskać jak najlepszy model użyliśmy jednak one-hot encodingu, ze względu na to, że jest to lepsza wersja tego co sami robimy w naszym kodowaniu.

```

[14]: df.info()

```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 435 entries, 0 to 434
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   handicapped_infants                  435 non-null    object
1   water_project_cost_sharing          435 non-null    object
2   adoption_of_the_budget_resolution  435 non-null    object
3   physician_fee_freeze                435 non-null    object
4   el_salvador_aid                     435 non-null    object
5   religious_groups_in_schools         435 non-null    object
6   anti_satellite_test_ban             435 non-null    object
7   aid_to_nicaraguan_contras          435 non-null    object
8   mx_missile                          435 non-null    object
9   immigration                         435 non-null    object
10  synfuels_corporation_cutback         435 non-null    object
11  education_spending                  435 non-null    object
12  superfund_right_to_sue              435 non-null    object
13  crime                               435 non-null    object
14  duty_free_exports                   435 non-null    object
15  export_administration_act_south_africa 435 non-null    object
16  political_party                     435 non-null    object
dtypes: object(17)
memory usage: 57.9+ KB

```

Modele działać, będą na dwóch ramkach, jedna cała, bez selekcji zmiennych i druga z odrzuceniem tych które są w bardzo niskim stopniu skorelowane z targetem.

```

[15]: df_dropped = df.
      ↪drop(['water_project_cost_sharing', 'immigration', 'export_administration_act_south_africa'],
      ↪axis=1)
X_dropped = df_dropped.iloc[:, :12]
y_dropped = df_dropped.iloc[:, 13]

X = df.iloc[:, :15]
y = df.iloc[:, 16]

```

```

[16]: x_train, x_test, y_train, y_test = train_test_split(df_dropped.iloc[:, :12],
      ↪df_dropped.iloc[:, 13], random_state=43)

```

```

[17]: import category_encoders as ce
      from sklearn.ensemble import GradientBoostingClassifier
      from sklearn.pipeline import Pipeline

      one_hot_encoder = ce.OneHotEncoder()

      one_hot = one_hot_encoder.fit_transform(X,y)

```

```
one_hot_encoder_dropped = ce.OneHotEncoder()
one_hot_dropped = one_hot_encoder_dropped.fit_transform(X_dropped,y_dropped)
```

7 GradientBoostingClassifier

```
[18]: clf = GradientBoostingClassifier(n_estimators=1000, learning_rate=.05,
    max_depth=3, random_state=997)

pipe_one_hot = Pipeline(
[
    ('transformer_one_hot', one_hot_encoder),
    ('classifier', clf)
])

pipe_one_hot_dropped = Pipeline(
[
    ('transformer_one_hot', one_hot_encoder_dropped),
    ('classifier', clf)
])

[19]: print(np.mean(cross_validate(pipe_one_hot, X, y, cv=11, scoring='accuracy').
    ↪get('test_score')))
print(np.mean(cross_validate(pipe_one_hot_dropped, X_dropped, y_dropped, cv=11,
    ↪scoring='accuracy').get('test_score')))
```

```
0.9493589743589744
0.9491841491841492
```

8 XGBClassifier

```
[20]: model=XGBClassifier(random_state=1,
    learning_rate=0.01, # Szybkość "uczenia" się
    booster='gbtree', # Jaki model wykorzystujemy (drzewo -
    ↪gbtree, liniowe - gblinear)
    nround = 1000, # Ilość iteracji boostingowych
    max_depth=3, # Maksymalna głębokość drzewa
    verbosity = 0
)

XGB_one_hot = Pipeline(
[
    ('transformer_one_hot', one_hot_encoder),
    ('classifier', model)
])

XGB_one_hot_dropped = Pipeline(
```

```
[
    ('transformer_one_hot', one_hot_encoder_dropped),
    ('classifier', model)
])
```

```
[21]: print(np.mean(cross_validate(XGB_one_hot, X, y, cv=11, scoring='accuracy').
    ↳get('test_score'))))
print(np.mean(cross_validate(XGB_one_hot_dropped, X_dropped, y_dropped, cv=11,
    ↳scoring='accuracy').get('test_score'))))
```

0.9515734265734267

0.9561188811188811

9 AdaBoostClassifier

```
[22]: model = AdaBoostClassifier(random_state=1, learning_rate=0.1, algorithm = 'SAMME.
    ↳R')
Ada_one_hot = Pipeline(
[
    ('transformer_one_hot', one_hot_encoder),
    ('classifier', model)
])

Ada_one_hot_dropped = Pipeline(
[
    ('transformer_one_hot', one_hot_encoder_dropped),
    ('classifier', model)
])
```

```
[23]: print(np.mean(cross_validate(Ada_one_hot, X, y, cv=7, scoring='accuracy').
    ↳get('test_score'))))
print(np.mean(cross_validate(Ada_one_hot_dropped, X_dropped, y_dropped, cv=7,
    ↳scoring='accuracy').get('test_score'))))
```

0.951649476995099

0.9562577719259747

9.1 Uwaga!

Ze względu na istotę problemu, czyli wybór między dwoma ugrupowaniami politycznymi, jedyną interesującą nas miarą jest accuracy, gdyż dla nas jest to bez różnicy czy pomylimy się mówiąc, że republikanin jest demokratą, czy też na odwrót.

10 TO DO:

Zamierzamy teraz skupić się na dwóch najlepiej zapowiadających się modelach, czyli AdaBoostClassifier i XGBClassifier. Chcemy poprawić ich parametry w taki sposób, aby otrzymać jak najlepsze

wyniki.